

Project Database

Due: December 16, 2016

1 Introduction

The Marine Life Institute, an aquarium and fish hospital in California, has a problem with their fish getting lost. Just last week a blue tang fish was swimming through the pipes, a whale shark leaped out of her exhibit into the ocean, and an octopus hijacked one of the Institute's trucks and released the fish it was carrying. The aquarium realizes that their internal database needs a huge overhaul to sort through this mess and put all the fish back where they belong. Currently to find the IDs and health information for the organisms at the Institute, employees have to look through the database on a computer in the facility's administrative office. Ideally, everyone would be able to search this data from any of the Institute's exhibits. Unfortunately, the database does not support multiple clients. Furthermore, even if it did, it was never written to be thread safe. The task to upgrade the aquarium's database has fallen to you.

In this project you will be creating a simple server to manage a Marine Life Database of key value pairs over a network. Aquarium specialists should be able to search for fish in the database, add new entries, and remove existing entries.

2 Support Code

The support code for this project consists of the following files:

- *server.c*: A C file containing the server program's main function and related code.
- *client.c*: A C file that produces the *client* program.
- *db.h*: A header containing function declarations for the database.
- *db.c*: A C file implementing database functionality.
- *comm.h*: A header defining abstraction for server side communication with clients.
- *comm.c*: Implementation of communication functions.
- *scripts/*: A directory containing database test scripts. Each of these scripts is just a series of client commands in a file, one per line. Feel free to make your own in addition to those provided.
- *database.pdf*: This handout.

In addition, you will need to write a Makefile for the project, which produces the executables *server* and *client*. You will only ever need to directly run *server* to start the server and run *client* to connect to the server.

Further, OSX may not implement some pthread functionality such as read/write locks and barriers. To accommodate this, we have attached the following files:

- *pthread_OSX.c*: A C file implementing missing OSX pthread functionality.
- *pthread_OSX.h*: A header file defining declarations and structs for OSX pthread functionality.

You can install this project by running

```
cs033_install database
```

The code contained in the handout produces a single-threaded, server-only version of the database. You will be modifying this to allow for multiple clients, a thread-safe database, and client cancellation.

You will only need to modify *server.c*, *db.c*, *comm.c*, and possibly *db.h* to complete this project.

2.1 Client Interface

Once you have a working server, users will be able to connect to it and interface with the database using the provided client. Your server will need to be able to handle several clients at a time, processing their commands as they come in. These clients will support only the client commands described below, while the *server* REPL only supports select server-only commands. The client interface receives queries from a script file or a REPL and sends them to the server. The server processes each query and sends an appropriate response to the client.

The database accepts the following commands:

- **a** *<key>* *<value>*: Adds *<key>* into the database with value *<value>*.
- **q** *<key>*: Retrieves the value stored with key *<key>*.
- **d** *<key>*: Deletes the given key and its associated value from the database.
- **f** *<file>*: Executes the sequence of commands contained in the specified file.

A client will close its connection to the server and exit upon reaching the end of a script file, or on CTRL+D if reading from the command line.

2.2 Testing Resources

In the *scripts/* directory, you are provided with a pair of files called *names2013.txt* and *names1880.txt* that you can use to quickly initialize the database. These files contain series of add commands to insert (name, frequency) pairs for babies born in the indicated years into the database. In addition, the *scripts* directory contains *test1.txt*, which you can use to test your database's thread safety in Part 2.

You can check that the tree output by the **p** command is correct with the following script:

```
cs033_db_check <txtfile>
```

where **txtfile** is the file containing the output of the **p** command.

To visualize the tree output by the **p** command, we have provided a script that creates a PNG image of the database. To use this script, enter the following command in a terminal:

```
cs033_db_vis <txtfile> <pngfile>
```

where `txtfile` is the file containing the output of the `p` command, and `pngfile` is the filename for the resulting PNG image.

3 Database

The database, implemented in *db.c*, consists of a collection of nodes organized in a binary search tree (which is not necessarily balanced). Each node contains a pointer to a left and right child, either or both of which may be null pointers. The key associated to a given node is lexicographically greater than all nodes in its left subtree, and lexicographically less than all nodes in its right subtree. (In other words, an in-order traversal of the tree nodes yields a lexicographical ordering of the corresponding keys.)

The database supports the following functions:

3.1 db_add()

The `db_add()` function calls `search()` to determine if the given key is already in the database. If the key is not in the database, the function creates a new node with the given key and value and inserts this node into the database as a child of the parent node returned by `search()`.

3.2 db_query()

The `db_query()` function calls `search()` to retrieve the node associated with the given key. If such a node is found, the function retrieves the value stored in that node and returns it.

3.3 db_remove()

The `db_remove()` function calls `search()` to retrieve the node associated with the given key. If such a node is found, the function must delete it while preserving the tree ordering constraints. There are three cases that may occur, depending on the children of the node to be removed:

- Both children are NULL: In this case, the function simply deletes the node and sets the corresponding pointer of the parent node to NULL.
- One child is NULL: In this case, the function replaces the node with its non-NULL child.
- Neither child is NULL: In this case, the function finds the leftmost child of the node's right subtree and replaces the removed node with this one. Since the replacement node has no left child, it is easy to remove it from its current position, and since it is the leftmost child of its subtree it can occupy the position of the deleted node and satisfy the tree's ordering constraints.

3.4 db_print()

The `db_print()` function performs a pre-order traversal of the tree, printing the node's representation and then recursively printing its left and right subtrees. It will attempt to print to a file with the given filename, or stdout if none is provided. Be sure to remove trailing whitespace (e.g. newlines) from the filename!

3.5 db_cleanup()

The `db_cleanup()` function **fre**es all dynamically-allocated nodes in the database. You should call this function before exiting on the server, and *only* when you are certain that no other threads are using or will use the database. You should use the variables in the `server_control_t` struct located near the top of `server.c` to ensure that all threads are terminated *before* you call `db_cleanup` in your main thread.

3.6 interpret_command()

The `interpret_command()` function gets called by the server to interpret the command, call database functions and store the response.

4 Assignment

This assignment is split up into three parts. We recommend starting this project early because you'll be applying a lot of different concepts. A completely functioning database will adhere to the following specifications:

- Part 1
 - Your server must be able to handle multiple clients. Each connection from a client to the server should be serviced by a thread. You must maintain a list of these client connection threads.
- Part 2
 - The database must implement fine-grained locking (see below for explanation).
 - Your server must be able to handle “s” (stop), “g” (go) commands using a condition variable and mutex, and “p” (print).
- Part 3
 - When the server receives an EOF from `stdin`, all client connection threads should be immediately terminated (via cancellation), after which the server should exit cleanly. You must also ensure that no new `client_t` structs are added to the thread list after you call `delete_all()` from the main thread when your server REPL terminates. Allowing this would cause your server to terminate with a non-empty thread list which would result in a memory leak! When debugging this you may find it useful to **assert** that the list is empty before calling `db_cleanup()` (Note, when a client loses connection to the server, it should shut itself down).

- When the server receives a SIGINT, all client connections should be immediately terminated via cancellation, after which the server should continue its input loop. Your implementation must allow users to spawn new clients after a SIGINT is delivered to the server. The server should behave exactly as it did before receiving a SIGINT (only now the database is potentially non-empty because of operations requested by previous clients).

5 Part 1

5.1 Networking

Your first task is to establish a TCP connection between a client and the server. Similar to the Networking Lab, you will be implementing a server using sockets.¹

To establish the connection, the client must find the server's internet address using `getaddrinfo()` and then set up a TCP connection with it. After finding the appropriate socket, the client creates an I/O stream to communicate with the server. Finally, the client uses `fputs()`, `fgets()` and `fflush()` to transmit and receive data.

Your next task is to implement the server functions (we have provided stubs) to handle multiple client connections. All clients should access the same, shared database.²

The server will be set up in a similar manner to the client, except it will also need to bind to the socket associated with the given port and accept connections from new clients. To accept connections, the server should have a thread that listens for new clients on the bound socket using `accept()`.³ Once `accept()` returns, the server will create an I/O stream and a new thread to handle the client connection. Take a look at the demo server provided with the first Network Programming lecture for a refresher if you need one!

Lastly, the server will use `comm_serve()` to read from and write to the client.

5.2 Multithreading

To implement the multithreaded version of the server, we suggest that you start with `client_constructor()` so that it spawns a new thread which executes `run_client()`. You will also need to ensure that `client_destructor()` is called at some point for each client, probably at the end of `run_client()`.

The server thread should stop accepting input on EOF. The easiest way to do this is probably to detach each client thread. Be careful when you clean up at the end – the database should be deleted, but not before all clients are done with it. (Hint: One way to achieve this is to maintain a thread-safe counter of the number of active threads.)

¹We are providing a fully functional *client.c*, but you are welcome to dissect and tamper with it. Keep in mind, however, that if you do, it will become more difficult for us to help you with bugs!

²Note that the database will not yet be thread-safe, so your program may behave incorrectly or crash if it receives input from more than one client at a time. You will fix this shortly.

³`accept()` blocks until the server receives a connection request.

6 Part 2

6.1 Thread Safety

Now that you have multiple threads, you must modify the database so that it is thread-safe. There are two principal ways to do this: apply *coarse-grained locking* and apply *fine-grained locking*. Both techniques are discussed below, but you must implement fine-grained locking.

In addition to making the database thread-safe, you will add some additional features to the server to facilitate testing.

6.1.1 Coarse-Grained Locking

The simplest way to ensure thread safety is to put a read/write lock on the whole database. Each thread should obtain an appropriate type of lock before accessing the database.

We recommend you implement and test coarse-grained locking first to get used to read/write locks before attempting the more difficult fine-grained locking scheme described below.

6.1.2 Fine-Grained Locking

Coarse-grained locking is easy to implement, but it is not very efficient. For any modifications to occur, a single thread must obtain exclusive access to the entire database. This strategy ignores the fact that nodes in the tree have some level of independence. A more efficient design would use *fine-grained locking*. In this design, each node in the tree has its own read/write lock. These locks must be carefully managed to maintain database consistency while allowing multiple threads to access and modify the database simultaneously.

To implement fine-grained locking, you should modify the `search()` function in `db.c` to accept an additional parameter specifying read- or write-locking. This function should lock the root node and percolate down the tree, locking each node *before* releasing the lock on its parent. You will also have to modify the other database functions so that they handle locking appropriately. Be sure to update `db_print()` because it does not use `search()`. You must think carefully about the operations involved to avoid deadlocks and ensure that the database stays consistent.

6.1.3 Testing Features

To test thread-safety, you should add the following functionality to your server code. When the server encounters the line “s”, all client threads should temporarily stop handling input. The line “g” should resume activity. This feature should allow you to test your database’s thread safety by pausing activity, creating several clients, and then running them all at once. To implement this feature, you should use a “stopped” flag with a condition variable (and associated mutex).

When the server encounters “p”, it should call `db_print()` on the database. This will help you test by allowing you to see the state of the database. You may also provide a filename to `db_print()` if you want the output in a file instead of standard out.

7 Part 3

In the last part of this assignment, you will add signal handling and cancellation to your database. Specifically, you must update the database so it supports the following behavior:

- When the server process receives an `EOF` from `stdin`, all clients should be immediately terminated, after which the program should exit cleanly.
- When the server process receives a `SIGINT` signal, all existing clients should terminate. The program should continue to accept input and create connections to new clients when requested.
- When working with this part of the assignment, it's important to ensure that your program is not leaking any memory support client thread data. You should ensure that the client thread data is cleaned up appropriately and that your thread list is completely empty before the server cleans up the database upon termination. To determine if a leak detected by `valgrind` is occurring due to this library behavior, check the trace of the reported leak; if the trace begins at `pthread_cancel()` or `pthread_cancel_init()`, and only references lines of code in library C files, then it's not a result of your code. You may also notice it if running a variable number of clients against your server produces the same roughly constant leak. This apparent leak is visible in the demo we have provided. If you are still unsure about a leak in your code, please talk to a TA or post on Piazza.

To implement these two features, you should maintain a list of all client connection threads. When a thread is created, it should add itself to the list. To terminate all current threads, we cancel all threads on the list. When a thread terminates (either from cancellation or naturally), it removes itself from the list.

Note that you must be careful about the timing involved — if the listener thread creates a new client, then the main (REPL) thread reaches an `EOF`, it may issue a “cancel all” command before the new client connection has added itself to the list. To address this, you should have some thread-safe mechanism for the server to indicate that it is no longer accepting new clients, and each new client should ensure that the server is still accepting clients before it adds itself to the list of clients and begins to serve the client with which it's associated.

7.1 Cancellation

To implement Part 2, you will need to use *cancellation*, a feature of the `pthread` library that allows for semi-asynchronous thread termination. A thread may be marked for cancellation at any time using the `pthread_cancel()` function; however, it is not cancelled immediately. When a thread marked for cancellation reaches a *cancellation point*, one of a set of library functions specified by the POSIX standard, it is terminated. You can find a list of POSIX-specified cancellation points using `man 7 pthreads`. Of the support code functions, `get_command()` acts as a cancellation point and `send_response()` may act as a cancellation point in some cases. You should not add your own cancellation points. It is safe to call `pthread_cancel()` on a thread any number of times until the thread terminates.

Each thread maintains a stack of *cleanup handlers* using the functions `pthread_cleanup_push()` and `pthread_cleanup_pop()`. When a cancelled thread reaches a cancellation point, the current cleanup handlers are executed in first-in-last-out order.

You will need to modify the client-handling threads to properly handle cancellation by installing appropriate cleanup handlers and modifying the cancel state as appropriate. For instance, the following block of code:

```
void foo(pthread_mutex_t my_mutex, pthread_cond_t my_cond) {
    pthread_mutex_lock(&my_mutex);

    while(!(some condition))
        // Cancellation point - the thread could stop executing here
        pthread_cond_wait(&my_cond, &my_mutex);

    // If the thread was cancelled in the pthread_cond_wait, this mutex would
    // never be unlocked
    pthread_mutex_unlock(&my_mutex);
}
```

could be replaced with:

```
// Cleanup handler called when thread is cancelled to unlock mutex so the thread
// doesn't quit while holding the lock
void cleanup_pthread_mutex_unlock(void *arg) {
    pthread_mutex_unlock((pthread_mutex_t *) arg);
}

void foo(pthread_mutex_t my_mutex, pthread_cond_t my_cond) {
    pthread_mutex_lock(&my_mutex);

    // Push the mutex unlock handler onto the stack of cleanup functions in case
    // the thread is cancelled while holding the lock
    pthread_cleanup_push(&cleanup_pthread_mutex_unlock, (void *)&my_mutex);

    while(!(some condition))
        pthread_cond_wait(&my_cond, &my_mutex); // Cancellation point

    // Pop and execute cleanup handler to release the lock regardless of whether
    // thread was cancelled
    pthread_cleanup_pop(1);
}
```

This way, the mutex will definitely be unlocked, even if the thread is cancelled while waiting.

Note that cancellation only occurs at most once per thread, so a cleanup handler that is executed as a result of explicit cancellation via `pthread_cancel()` may safely contain cancellation points. However, cancellation *can* occur while executing a cleanup handler that is explicitly called via `pthread_cleanup_pop()`.

Each thread has a *cancel state*, which sets whether or not a thread can be terminated at a cancellation point. The cancel state can be controlled with the `pthread_setcancelstate()` function. Consider when it may be necessary to call this function in your implementation.

Note: The `pthread` manual page states that `pthread_rwlock_{wr, rd}lock()` functions **may** be cancellation points. These functions **are not** implemented as cancellation points on Linux, therefore you may assume that they are not cancellation points in your implementation.

7.2 Signal Handling

Signals occur at the process level, making them somewhat difficult to handle in multithreaded code. For this project, you will handle signals in a separate thread that runs alongside the server and client threads. To set this up, you should first mask off `SIGINT` for the entire process using `pthread_sigmask()`. Then, create a special signal monitoring thread that uses the `sigwait()` function to listen for `SIGINT` signals and cancel running clients.

8 Demos

We have provided a demo binary for this project: `cs033_db_demo_server`. This is a fully-functional DB server, with all parts of the assignment implemented, for you to get a better feel for the finished result.

9 GDB Tips

As you have probably seen, debugging programs involving multiple processes can be much more confusing than debugging a single component on its own. However, GDB has several tools that may be able to come to the rescue!

A basic functionality of GDB in general is to pause program execution at a specified point and investigate the state of the system. The best way to do that here is to view the current threads with `info threads`. The current thread (GDB is able to watch the stack of a single thread of execution at a time) will have a star next it, and you can then switch to a thread with `thread <thread number>`. Each thread can give you a `backtrace`, and if you want to see them all at once, you can try something like `thread apply all backtrace`, which will give you the backtrace for all threads. This works for other commands as well!

If you are facing a deadlock, you will want to interrupt execution with `CTRL+Z` and dig your way through mutexes and so forth to identify which resources are causing the problem. You can move through a trace with `frame <frame number>`. If you find that a particular thread is stuck in, say, `pthread_mutex_lock()`, you'll be able to see which mutex is being waited on, and which thread owns it. Then switch to that thread and see what it is waiting for, and so on. For more on debugging deadlocks, visit this helpful link⁴.

Keep in mind other common debugging techniques like setting watchpoints, calling functions within gdb, setting conditional breakpoints, etc. and you'll have a much better time!

⁴https://en.wikibooks.org/wiki/Linux_Applications_Debugging_Techniques/Deadlocks

10 Handin

To hand in your database implementation, run

```
cs033_handin database
```

from your project working directory. Make sure you include all C files, your Makefile, and a README. If you wish to change your handin, you can do so by re-running the handin script. Only your most recent handin will be graded.

11 Grading

This project will be graded according to the following categories, ordered by weight:

- *Functionality.*
 - Your client and server should communicate correctly.
 - You should cleanly start and close connections.
 - You should correctly implement networking and fine-grained locking.
- *Correctness.*
 - Your server and database should be able to correctly handle multiple clients at once.
 - Your database should be memory safe and should work.
- *Style.* Your code should be readable and commented and have a README.